



# Azure AI Services Infrastructure

Enterprise-Ready AI Platform with Network Isolation & Private Endpoints



## Virtual Network: ai-platform-vnet-region (192.168.0.0/16)



### agent-subnet

192.168.0.0/24

Delegated to Microsoft.App/environments



AI Agent Runtime Environment



Network Injection Target



Capability Host Deployment



### pe-subnet

192.168.1.0/24

Private Endpoints Subnet

Network isolation for all services

DNS resolution hub



### Private Endpoints (in pe-subnet)



AI Services

192.168.1.8-10



Search Service

192.168.1.6



Cosmos DB

192.168.1.4-5



Storage Account

192.168.1.7



### Private DNS Zones

- privatelink.cognitiveservices.azure.com
- privatelink.openai.azure.com
- privatelink.services.ai.azure.com

- privatelink.search.windows.net
- privatelink.documents.azure.com
- privatelink.blob.core.windows.net





## Azure AI Services Hub

ai-services-platform



### AI Foundry Hub Service:



Project: ai-platform-project



Model: GPT-4o (GlobalStandard)



Capability Host: Agents

#### Connected Resources:

- Vector Store: ai-search-platform
- Storage: ai-storage-platform
- State Store: ai-cosmos-platform
- Network Injection: agent-subnet



#### External Connections:

- Microsoft Fabric Workspace
- Databricks Genie Analytics
- Power BI Premium Datasets
- Data Lake Gen2 Sources



#### Security Features:

- Content Safety (DefaultV2 Policy) • AAD Authentication • Private Endpoint Only • System Assigned Identity



## Azure AI Search

ai-search-platform

#### Configuration:

- Standard SKU (1 replica, 1 partition) • Vector Store for RAG Operations • System Assigned Identity • AAD Authentication Enabled • Public Network Access: Disabled • Semantic Search: Disabled



#### Network Security:

- Private Endpoint: 192.168.1.6 • DNS Zone: privatelink.search.windows.net • Connection State: Approved



## Azure Cosmos DB

ai-cosmos-platform

#### State Store Configuration:

- SQL API with Session Consistency • Database: platform\_memory • Autoscale: 100-1000 RU/s per container • TLS 1.2+ Only, Local Auth Disabled



#### State Store Containers:

- agent-entity-store (Agent persistence) • conversation-state-store (Conversation state) • system-state-store (System state)



#### Security:

- Role-based Access Control • Private Endpoint: 192.168.1.4-5 • Geo-redundant Backup



## Azure Storage Account

ai-storage-platform

#### Storage Configuration:

- StorageV2, Standard ZRS • Services: Blob, File, Queue, Table • Hot Access Tier • TLS 1.2+ Only



#### Containers:

- platform-blobstore (Project artifacts) • platform-agent (Agent runtime data)



#### Security Features:

- No Public Blob Access • Shared Key Access: Disabled • Private Endpoint: 192.168.1.7 • Cross-tenant Replication: Disabled

- AI/ML Services
- Search & Vector Store
- Database Services
- Storage Services
- Agent Runtime Subnet
- Private Endpoints Subnet
- Private Endpoints